



# Prometheus

The light, the shadow and the livers...



**Martin Chodúr**

FUSAKLA

# Project not a company



prometheus/prometheus



 fabxc



 juliusv




 brian-brazil



 beorn7

# Incubating in Cloud Native Computing Foundation



# Data model

- Multidimensional model
- Every time-series identified by name and labels
- Metric name is just specifically handled label
- Sample consists of millisecond precision timestamp and float64
- 4 basic metric types: Counter, Gauge, Summary, Histogram

# Counter

```
http_requests_total{status="200", endpoint="/metrics"} 10
```

- Always rising, cannot decrease at any point
- Should have the `_total` suffix in metric name
- Engine needs to deal with counter resets (application restart)
- Should be used with the `rate` , `irate` and `increase` functions

# Gauge

```
node_memory_usage_bytes{node="fusak1a.cz"} 23005456
```

- Represents actual value of some phenomenon
- Accuracy relies on the scrape interval you use
- Can't be used with the `rate` , `irate` and `increase` functions

# Histogram

```
http_request_duration_seconds_count{endpoint="/metrics"} 5
http_request_duration_seconds_sum{endpoint="/metrics"} 1
http_request_duration_seconds_bucket{endpoint="/metrics", le="0.1"} 1
http_request_duration_seconds_bucket{endpoint="/metrics", le="0.5"} 3
...
http_request_duration_seconds_bucket{endpoint="/metrics", le="+Inf"} 5
```

- Metric type consisting of set of Counters
- Describes count, total sum and distribution in given buckets
- Using the `histogram_quantile` it's possible to compute quantiles



# Summary

```
http_request_duration_seconds_count{endpoint="/metrics"} 5
http_request_duration_seconds_sum{endpoint="/metrics"} 1
http_request_duration_seconds{endpoint="/metrics", quantile="0.5"} 1
http_request_duration_seconds{endpoint="/metrics", quantile="0.9"} 3
...
http_request_duration_seconds{endpoint="/metrics", quantile="0.99"} 5
```

- Same as Histogram only evaluated on exporter side
- No need to compute the quantile on the Prometheus side
- Hides bucket distribution (Histogram is more used)

# Gathering the data

- Prometheus uses pull based model
- Application exposes metrics on HTTP endpoint
- Periodically scrapes data from targets
- Variety of client libraries
- Exporters. A lots of them...

# Exposition format

```
# HELP grafana_info Information about the Grafana
# TYPE grafana_info gauge
grafana_info{version="5.1.3"} 1
# HELP grafana_instance_start_total counter for started instances
# TYPE grafana_instance_start_total counter
grafana_instance_start_total 1
# HELP grafana_page_response_status_total page http response status
# TYPE grafana_page_response_status_total counter
grafana_page_response_status_total{code="200"} 11808
grafana_page_response_status_total{code="unknown"} 89
# HELP grafana_proxy_response_status_total proxy http response status
# TYPE grafana_proxy_response_status_total counter
grafana_proxy_response_status_total{code="200"} 981338
grafana_proxy_response_status_total{code="400"} 3
grafana_proxy_response_status_total{code="500"} 49672
```

# Open Metrics

OpenMetrics is a working group to determine a standard for exposing metrics data, influenced by the Prometheus exposition format.



RichiH/OpenMetrics


# Running and operating it

- Prometheus is single binary with no dependencies written in Go
- Configured by YAML file
- Prometheus does NOT have any configuration features...

That leaves to the operator

- Can be CPU and memory intensive
- Needs persistent filesystem (do not use NFS!)

# Demo time



# Go instrumentation

```
var httpRequestDurationHistogram = prometheus.NewHistogramVec(  
    prometheus.HistogramOpts{  
        Name: "http_request_duration_seconds",  
        Help: "HTTP request latency distribution",  
    },  
    []string{"status", "endpoint"},  
)  
  
prometheus.MustRegister(httpRequestDurationHistogram)  
  
httpRequestDurationHistogram.WithLabelValues("200", "/demo").Observe(duration.Seconds())  
  
http.Handle("/demo/metrics", promhttp.Handler())
```


# Prometheus configuration

```
global:
  scrape_interval:    5s
  evaluation_interval: 10s


scrape_configs:
  - job_name: 'prometheus'
    static_configs:
      - targets: ['localhost:9090']
  - job_name: 'error-server'
    metrics_path: "/demo/metrics"
    scheme: https
    static_configs:
      - targets: ['fusakla.cz']
```



# TSDB

- Prometheus uses own database  [prometheus/tsdb](https://github.com/prometheus/tsdb)
- Uses 2h WAL and then stores data in blocks on filesystem
- Blocks are compacted to longer blocks (6h, 12h)
- Not meant for longer retentions
- Prometheus allows backups using API

# Service discovery

- Main feature for pull based model
- Implements cca 12 different SD allowing dynamic target load
- Well integrated with Kubernetes  [coreos/prometheus-operator](https://github.com/coreos/prometheus-operator)

# PromQL

- Powerful query language with variety of operators and functions
- Results in one of these: Instant vector, Range vector, Scalar, String
- Allows aggregating on all the labels
- Cannot create Range vector for function result
- Turing complete

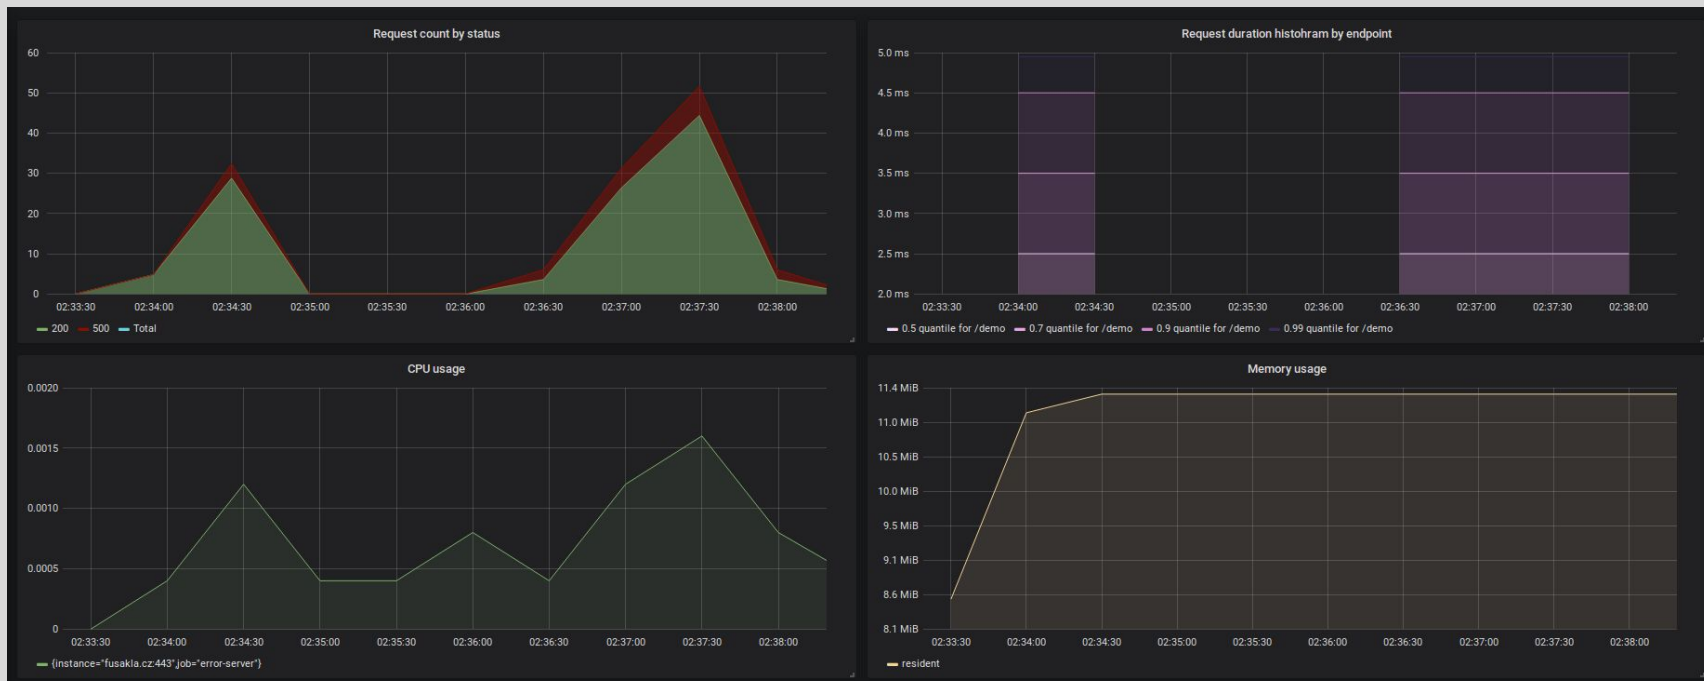
```
rate(http_requests_total{job="api-server"}[5m])
```

# Data and PromQL

```
increase(http_requests_total{endpoint="/metrics"}[1m])
```

	2	4	5	6	7	8	15	1	2	4	5	5	7	8		
Raw data		-	+	-	+	-	+	-	+	-	+	-	+	-		
		1m														
Range vector		-	-	-	+	-	-	-	+	-	-	-	+	-		
		(4-2)		(6-5)		--		(15-8)		(2-1)		(5-4)		(7-5)		(8-7)
		/60		/60				/60		/60		/60		/60		/60
Result		0.03		0.01		--		0.1		0.01		0.01		0.03		0.01

# Grafana dashboard



# What's next

- Using recording rules to avoid repetitive query evaluation
- Explore Prometheus HTTP api and remote read and write API
- HA Alerting using alert rules and Alertmanager
- Hierarchical federation
- Relabeling configs
- Long term storage

# Closing remarks

Every time when Prometheus exits, last thing it does is writing to the log:

```
level=info caller=main.go:599 msg="See you next time!"
```

Wait when you mess up configuration few times in a row... spiteful little troll



# Ask questions to

Me:



FUSAKLA



m.chodur@seznam.cz



Martin Chodur

Community:



IRC #prometheus



prometheus-users



@PrometheusIO

...it's always better to waste someone else's time than yours